

Lab2 Report

Multi-Layer Perceptron

Eduardo Delgado Coloma Bier (s3065979)
Mikel Orbea Sotil (s3075001)

March 3, 2016

1 Theory Questions

- a) The credit assignment problem is the process of determining which weights are responsible for the output of the network. In other words, which weights are to blame and need to be changed when the output is not what we want? Or, when things go well, which weights are responsible and shouldn't be changed?
- b) The error in a hidden layer is determined by the output of the hidden layer and the goal.
- c) The sigmoid function is defined by:

$$\sigma(a) = \frac{1}{1 + \exp[\frac{-(a-\theta)}{\rho}]} \quad (1)$$

θ and ρ are constants, while a is the variable given by the activation (summed weights \times input). While θ is responsible for when the curve starts going up (the larger the θ , the longer it takes), ρ is responsible for how smoothly it does so. With a low ρ one can expect a very steep curve. The derivative of the sigmoid function is:

$$\sigma'(a) = \frac{1}{\rho} \cdot \sigma \cdot (1 - \sigma) \quad (2)$$

- d) If we initialize the weights with a very high value, chances are our activation will have a high value as well. This means that we would be on the right part of the sigmoid function, where the derivative is close to 0, which is something we don't want because the weights would be adjusted really slowly.
- e) One criteria we can use to choose when to stop learning is choosing a minimum error ϵ where we would stop learning if the error is smaller than ϵ . However, choosing the ϵ could be tricky and completely arbitrary. Another possibility would be to stop as soon as the error starts getting bigger. This method, however is not ideal since if we reach a local minimum of the function, we'd probably stop there instead of the actual optimal solution. A third stop criteria would be to keep track of the error and if it's never lower than the last minimal error you found for a certain (big) number of epochs, then you rollback to that minimal error you found. Choosing the amount of epochs, though, can also be tricky and lead to sub optimal solutions.
- f) Aside from changing the learning rate, one can speed the learning of a network by choosing appropriate initial weights, neither too high or too low. Both too high and too low values for the weights would lead to σ' values close to zero, which in turn would make every step taken by the network really small.
- g) To verify that the network is generalizing for a set of training data we use cross-validation. This method consists on setting aside part of the training set to be used as a validation set after the training. If the error in the validation set starts increasing, then the network is starting to overfit and should be stopped. A problem with this method is that sometimes the error in the validation set only grows temporarily.
- h) Overfitting happens when a network trains for too long with a set of data and ends up incorporating too many details from that learning set. When that happens, the network works really well for the learning set, but when different inputs (not in the training set) are given to it, they are not well classified.
- i) Network pruning is the deletion of nodes within the network that are considered not important for the network to work. The network is therefore smaller and simpler, resulting in a better performance overall. The basic idea of pruning method is: start with a large enough network, train it with a set of data, determine the importance of the weights, remove the least important weight, retrain the network and repeat that until you get a reasonably sized network. The

difference in each pruning method is basically choosing when and which node to cut. One way of doing it is always removing the weight with the lowest value. Another possibility is removing the connection with the smallest contribution (to the network) variance.

2 An MLP on paper

a) Neural network




Figure 1: Neural networks with 2 input neurons, 3 hidden neurons and 2 output neurons

- b) That is considered a two layer network because the input layer is not really considered a layer as it does no computation whatsoever.
- c) W^h is a 2×3 matrix, while W^o is a 3×2 matrix.
- d) The weights on the same column all go to the same neuron.
- e) Colored weights




Figure 2: Neural networks with highlighted weights: w_{12}^h in red and w_{21}^o in blue

f) With the augmented input layer, W^h is a 3×3 matrix, while W^o is a 3×2 matrix.




Figure 3: Neural network with augmented input layer

g) The values of this table can be found on the first row of W^h

input neuron	weights to first hidden neuron
1	0.3
2	0.6
3	0.5

Table 1: Connections from the augmented input neurons to the first hidden neuron

h) Given the input vector $\vec{x} = \begin{pmatrix} 1 \\ 0 \\ -1 \end{pmatrix}^T$ and the weight vector $\vec{w} = \begin{pmatrix} 0.3 \\ 0.6 \\ 0.5 \end{pmatrix}$, we can easily calculate the activation using the following equation:

$$a = \vec{x} \cdot \vec{w} = 1 \cdot 0.3 + 0 \cdot 0.6 + (-1) \cdot 0.5 = -0.2 \quad (3)$$

i) By multiplying $\vec{x} \cdot W^h$ we end up with an array \vec{a} whose i -th is exactly $\vec{x} \cdot W^h_i$, where W^h_i is the i -th column of the W^h matrix. This means that the array calculated by the multiplication holds the activation for each neuron.

j) To calculate the activation of the output layer, we simply need to multiply the output vector of the hidden layer y^h , a 1×3 array, by the weights of the output layer, W^o , a 3×2 matrix, similar to what we did to calculate the activation of the hidden layer. Therefore, we get:

$$\vec{a}^o = y^h \cdot W^o \quad (4)$$

Notice that the dimensions match and that, as expected, we end up with 2 outputs.

3 Implementing an MLP in Matlab

1. MLP

```
1 % mlp.m Implementation of the Multi-Layer Perceptron
2
3 clear all
4 close all
5
6 examples = [0 0;1 0;0 1;1 1];
7 goal = [0.01 0.99 0.99 0.01]';
8
9 min_error = 0.01;
10
11 % Boolean for plotting the animation
12 plot_animation = true;
13
14 % Parameters for the network
15 learn_rate = 0.2;           % learning rate
16 max_epoch = 5000;           % maximum number of epochs
17
18 mean_weight = 0;
19 weight_spread = 5;
20
21 n_input = size(examples,2);
22 n_hidden = 20;
23 n_output = size(goal,2);
24
25 % Noise level at the input
26 noise_level = 0.05;
27
28 % Activation of the bias node
29 bias_value = -1;
30
31
32 % Initializing the weights
33 w_hidden = rand(n_input + 1, n_hidden) .* weight_spread - weight_spread/2 +
    mean_weight;
34 w_output = rand(n_hidden, n_output) .* weight_spread - weight_spread/2 +
    mean_weight;
35
36 % Start training
37 stop_criterium = 0;
38 epoch = 0;
39
40 while ~stop_criterium
41     epoch = epoch + 1;
42
43     % Add noise to the input data.
44     noise = randn(size(examples)) .* noise_level;
45     input_data = examples + noise;
46
47     % Append bias to input data
48     input_data(:,n_input+1) = ones(size(examples,1),1) .* bias_value;
```

```

49
50 epoch_error = 0;
51 epoch_delta_hidden = 0;
52 epoch_delta_output = 0;
53
54 % FROM HEREON YOU NEED TO MODIFY THE CODE!
55 for pattern = 1:size(input_data,1)
56
57     % Compute the activation in the hidden layer
58     hidden_activation = input_data(pattern, :) * w_hidden;
59
60     % Compute the output of the hidden layer (don't modify this)
61     hidden_output = sigmoid(hidden_activation);
62
63     % Compute the activation of the output neurons
64     output_activation = hidden_output * w_output;
65
66     % Compute the output
67     output = output_function(output_activation);
68
69     % Compute the error on the output
70     err = goal(pattern, :) - output;
71     output_error = 0.5 * err;
72
73     % Compute local gradient of output layer
74     local_gradient_output = d_sigmoid(output_activation) .* (err);
75
76     % Compute the error on the hidden layer (backpropagate)
77     hidden_error = 0.5 * (goal(pattern, :) - hidden_output);
78
79     % Compute local gradient of hidden layer
80     local_gradient_hidden = d_output_function(hidden_activation) .* (w_output *
81     local_gradient_output)';
82
83     % Compute the delta rule for the output
84     delta_output = learn_rate .* hidden_output' * local_gradient_output;
85
86     % Compute the delta rule for the hidden units;
87     delta_hidden = learn_rate .* input_data(pattern, :) *
88     local_gradient_hidden;
89
90     % Update the weight matrices
91     w_hidden = w_hidden + delta_hidden;
92     w_output = w_output + delta_output;
93
94     % Store data
95     epoch_error = epoch_error + (output_error).^2;
96     epoch_delta_output = epoch_delta_output + sum(sum(abs(delta_output)));
97     epoch_delta_hidden = epoch_delta_hidden + sum(sum(abs(delta_hidden)));
98 end
99
100 % Log data
101 h_error(epoch) = epoch_error / size(input_data,1);

```

```

100 log_delta_output(epoch) = epoch_delta_output;
101 log_delta_hidden(epoch) = epoch_delta_hidden;
102
103 % Check whether maximum number of epochs is reached
104 if epoch > max_epoch
105     stop_criterium = 1;
106 end
107
108 % Implement a stop criterion here
109 if epoch_error < min_error
110     stop_criterium = 1;
111 end
112
113 % Plot the animation
114 if and((mod(epoch,20)==0),(plot_animation))
115     emp_output = zeros(21,21);
116     figure(1)
117     for x1 = 1:21
118         for x2 = 1:21
119             hidden_act = sigmoid([(x1/20 - 0.05) (x2/20 -0.05) bias_value] *
w_hidden);
120             emp_output(x1,x2) = output_function(hidden_act * w_output);
121         end
122     end
123     surf(0:0.05:1,0:0.05:1,emp_output)
124     title(['Network epoch no: ' num2str(epoch)]);
125     xlabel('input 1: (0 to 1 step 0.05)')
126     ylabel('input 2: (0 to 1 step 0.05)')
127     zlabel('Output of network')
128     zlim([0 1])
129 end
130
131 end
132
133 % Plotting the error
134 figure(2)
135 plot(1:epoch,h_error)
136 title('Mean squared error vs epoch');
137 xlabel('Epoch no. ');
138 ylabel('MSE');
139
140 % Add additional plot functions here (optional)

```

Listing 1: mlp.m

2. Sigmoid Function

```

1 %this functions calculates the sigmoid
2 function [output] = sigmoid(x)
3     output = 1 ./ (1 + exp(-x));
4 end

```

Listing 2: sigmoid.m

3. Output Function

```
1 function [output] = output_function(x)
2     output = sigmoid(x);
3 end
```

Listing 3: outputfunction.m

4. Derivative of the Sigmoid Function

```
1 %this functions calculates the differential of the sigmoid
2 function [output] = d_sigmoid(x)
3     temp = sigmoid(x);
4     output = temp .* (1 - temp);
5 end
```

Listing 4: dsigmoid.m

5. Derivative of the Output Function

```
1 %this functions calculates the differential of the output function
2 function [output] = d_output_function(x)
3     output = d_sigmoid(x);
4 end
```

Listing 5: doutputfunction.m

6. MLP Sine

```
1 clear all
2 close all
3
4 % The number of examples taken from the function
5 n_examples = 5;
6
7 examples = (0:2*pi/(n_examples-1):2*pi)';
8 goal = sin(examples);
9
10 % Boolean for plotting animation
11 plot_animation = true;
12 plot_bigger_picture = false;
13
14 % Parameters for the network
15 learn_rate = 0.05;           % learning rate
16 max_epoch = 5000;           % maximum number of epochs
17
18
19 mean_weight = 0;
20 weight_spread = 1;
21
22 n_input = size(examples,2);
23 n_hidden = 20;
24 n_output = size(goal,2);
25
26 % Noise level at input
27 noise_level = 0.05;
28
```



```

29 bias_value = -1;
30
31 % Initializing the weights
32 w_hidden = rand(n_input + 1, n_hidden) .* weight_spread - weight_spread/2 +
    mean_weight;
33 w_output = rand(n_hidden, n_output) .* weight_spread - weight_spread/2 +
    mean_weight;
34
35 % Start training
36 stop_criterium = 0;
37 epoch = 0;
38
39 min_error = 0.01;
40
41 while ~stop_criterium
42     epoch = epoch + 1;
43
44     % Add noise to the input
45     noise = randn(size(examples)) .* noise_level;
46     input_data = examples + noise;
47
48     % Append bias
49     input_data(:,n_input+1) = ones(size(examples,1),1) .* bias_value;
50
51     epoch_error = 0;
52     epoch_delta_hidden = 0;
53     epoch_delta_output = 0;
54     for pattern = 1:size(input_data,1)
55
56         % Compute the activation in the hidden layer
57         hidden_activation = input_data(pattern, :) * w_hidden;
58
59         % Compute the output of the hidden layer (don't modify this)
60         hidden_output = sigmoid(hidden_activation);
61
62         % Compute the activation of the output neurons
63         output_activation = hidden_output * w_output;
64
65         % Compute the output
66         output = output_function(output_activation);
67
68         % Compute the error on the output
69         err = goal(pattern, :) - output;
70         output_error = 0.5 * err;
71
72         % Compute local gradient of output layer
73         local_gradient_output = err;
74
75         % Compute the error on the hidden layer (backpropagate)
76         hidden_error = 0.5 * (goal(pattern, :) - hidden_output);
77
78         % Compute local gradient of hidden layer
79         local_gradient_hidden = d_output_function(hidden_activation) .* (w_output *

```

```

    local_gradient_output')');
80
81     % Compute the delta rule for the output
82     delta_output = learn_rate .* hidden_output' * local_gradient_output;
83
84     % Compute the delta rule for the hidden units;
85     delta_hidden = learn_rate .* input_data(pattern, :) *
local_gradient_hidden;
86
87     % Update the weight matrices
88     w_hidden = w_hidden + delta_hidden;
89     w_output = w_output + delta_output;
90
91     % Store data
92     epoch_error = epoch_error + (output_error).^2;
93     epoch_delta_output = epoch_delta_output + sum(sum(abs(delta_output)));
94     epoch_delta_hidden = epoch_delta_hidden + sum(sum(abs(delta_hidden)));
95 end
96
97 h_error(epoch) = epoch_error / size(input_data,1);
98 log_delta_output(epoch) = epoch_delta_output;
99 log_delta_hidden(epoch) = epoch_delta_hidden;
100
101 if epoch > max_epoch
102     stop_criterium = 1;
103 end
104
105 % Add your stop criterion here
106 if epoch_error < min_error
107     stop_criterium = 1;
108 end
109
110 % Plot the animation
111 if and((mod(epoch,20)==0),(plot_animation))
112     %out = zeros(21,1);
113     nPoints = 100;
114     input = linspace(0, 2 * pi, nPoints);
115     for x=1:nPoints
116         h_out = sigmoid([input(x) bias_value] * w_hidden);
117         out(x) = output_function(h_out * w_output);
118     end
119     figure(1)
120     plot(input,out,'r-','DisplayName','Output network')
121     hold on
122     plot(input,sin(input),'b-','DisplayName','Goal (sin[x])')
123     hold on
124     scatter(examples, goal, 'DisplayName', 'Voorbeelden')
125     hold on
126     title(['Output and goal. Epoch: ' num2str(epoch)]);
127     xlim([0 2*pi])
128     ylim([-1.1 1.1])
129     set(gca,'XTick',0:pi/2:2*pi)
130     set(gca,'XTickLabel',{'0','1/2 pi','pi','3/2 pi ','2 pi'})

```

```

131         xlabel('Input')
132         ylabel('Output')
133         legend('location','NorthEast')
134         hold off
135     end
136
137 end
138
139
140 % Plot error
141 figure(2)
142 plot(1:epoch,h_error)
143 title('Mean squared error vs epoch');
144 xlabel('Epoch nr. ');
145 ylabel('MSE');
146
147 %Plot the bigger picture
148 if plot_bigger_picture
149     figure(3)
150     in_raw = (-5:0.1:15)';
151     in_raw = horzcat(in_raw,(bias_value*ones(size(in_raw))));
152     h_big = sigmoid(in_raw * w_hidden);
153     o_big = output_function(h_big * w_output);
154
155     plot(-5:0.1:15,o_big,'r-','DisplayName','Output network')
156     hold on
157     plot(-5:0.1:15,sin(-5:0.1:15),'b-','DisplayName','Goal (sin[x])')
158     hold on
159     scatter(examples, sin(examples), 'DisplayName', 'Examples');
160     hold off
161     xlabel('Input')
162     ylabel('Output')
163     legend('location','NorthEast')
164     title('The bigger picture')
165 end

```

Listing 6: mlpsinus.m

4 Testing the MLP

- a) It is not guaranteed that the network will always find a solution. That happens because the number of nodes in the hidden layer is too small, which causes the network to have too few weights to play around with in the hidden layer. This makes it so that the network sometimes doesn't converge to a solution.
- b) About 2500 epochs in average are needed to find a solution with these settings.
- c) The network is unable to find a solution because the noise level is now too high. The noise level makes it harder for the inputs to be correctly classified because the inputs end up being too altered to be identified.
- d) The number of epochs needed to reach the solution is greatly decreased (370, down from 2500), so the network is much more efficient. That happens because now the weights are not so low anymore, which causes the local gradient of the hidden layer to be larger. That impacts the

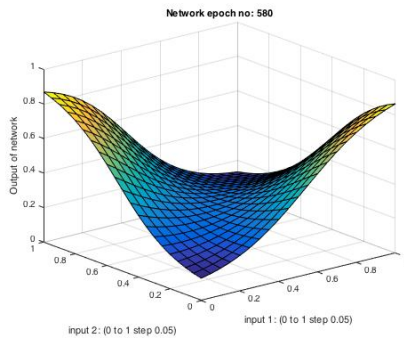


Figure 4: Solution 1

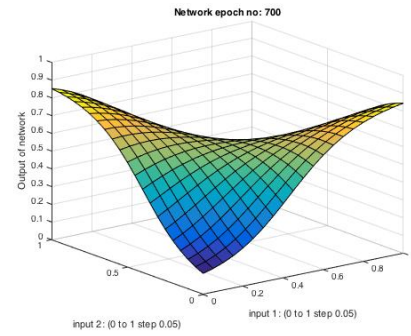


Figure 5: Solution 2

delta rule for the hidden layer directly, making it larger, which means larger steps (at least in the first few epochs). Therefore, the error decreases faster, and fewer epochs are needed to reach a reasonable weight.

- e) The two different solutions to the XOR problem are as shown above. Notice that the output of the center part of Solution 1 (Figure 4) is (close to) 0, while on Solution 2 (Figure 5) it is (close to) 1. That behavior is due to the fact that the points in the center, where x_1 and x_2 have similar values get undefined behavior for XOR, since the network doesn't know if they are representing a 0 or a 1. The network therefore has to make a "decision" regarding how they'll be treated.
- f) The graph of the error is usually a (noisy) decreasing exponential function. Since the delta rule is a linear function, the weights are changed linearly too. That means that the non-squared error also changes linearly, which makes the squared error decrease exponentially.

5 Another function

- a) Yes, although not perfectly, the neural network is able to learn the sine function, as shown in Figure 6

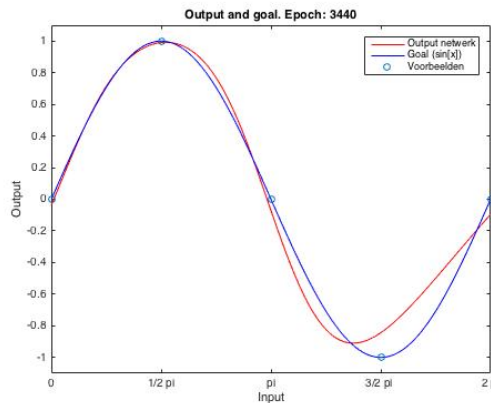


Figure 6: The network's learned function vs the sine function

- b) The network is still able to learn with fewer examples. This phenomenon corresponds to the fact that a neural network reaches a generalized solution that is independent of the input.

- c) The domain of the network is determined by the input set. If an input is outside of the domain, the network doesn't really know what to do with it, which leads to a wrong classification of that input. It's interesting to see that when the input is still close to the domain, the network is still able to give acceptable results, but as it gets further and further away from the domain, the results are terrible.
- d) At least 3 neurons are required to learn the sine function.
- e) The XOR learning network still works. The reason for that is that the sigmoid function, which was originally used as the output function for the XOR, does the same thing as the identity function in terms of making a continuous function that can be derived.